

# Homework 5 Introduction to Big Data Systems

Christoffer Brevik

October 27, 2024

# 1 Questions for MapReduce paper

Even though the questions ask me to find a solution to both problems, we actually see that both of these issues are solved the paper on MapReduce. I therefore focus on these solutions when answering the questions.

## Q1: Reducing Network Traffic during Shuffle Phase

In the MapReduce framework, the shuffle phase involves transferring the intermediate data generated by the map tasks to the reducers, which can result in significant network traffic. This is especially bad if there is significant repetition in the intermediate keys produced during the map phase. To address this issue, one effective approach is to implement a **combiner function**. The combiner function performs a local, partial aggregation of the intermediate data directly on the mapper nodes. This is very similar to the reduce operation, but it is handled during the map phase and the output is written to an intermediate file that will be sent to a reduce task. By utilizing this, the amount of data transferred over the network can be greatly minimized, thereby reducing network load.

## Q2: Addressing Slow Mapper or Reducer Tasks (Stragglers)

While data is typically distributed evenly, failures and bugs in code may cause certain sections to take a longer time to process during both the mapping and reduction phases. These slow tasks, known as *stragglers*, can delay the completion of a MapReduce job. The MapReduce framework mitigates this issue by employing a **backup task** mechanism.

When a job nears completion, the master node initiates duplicate (backup) instances of slow tasks. The job is completed as soon as either the original or any of the backup tasks finishes, after which any duplicate tasks are terminated. This approach helps ensure that the overall job execution is not delayed by any single slow task, thereby improving the completion time of the entire MapReduce operation.

In cases where specific bugs cause certain tasks to fail repeatedly, MapReduce provides an option to **skip problematic records**. If a map or reduce function crashes on a specific record multiple times, the framework detects this repeated failure and skips the problematic record in order to allow the job to make forward progress. This feature is particularly useful when dealing with data inconsistencies or errors that cannot be resolved at runtime, thus improving the reliability and robustness of the MapReduce job even in the presence of repeatable bugs.

## 2 Implementation of MapReduce

### How input data is structured

The `map` function processes each line of input data, which represents edges from a source node to a destination node. The data format, as shown in the ReadMe, is the following:

```
a 1 2 0
a 3 4 0
a 5 1 0
a 2 4 0
a 4 5 0
a 2 5 0
a 2 3 0
a 3 2 0
```

In this format, each line starts with the label `a`, followed by integers representing source and destination nodes. The last number is most likely the weight, and is not used for counting out-degrees.

### Updated `run_od.sh`

I modified the `run_od.sh` to automatically save the files after running, as well as renaming some variables as these still referred to variables as "words". The updated file is the following:

```
#!/bin/bash

OUTPUTPATH="od_res/"

echo ===== OutDegree JAVA VERSION =====

echo ===== Compile =====
javac -classpath '/hadoop/bin/yarn classpath' OutDegree.java
jar cf od.jar OutDegree*.class
echo

echo ===== Clear old output files on HDFS =====
/hadoop/bin/hdfs dfs -rm -r $OUTPUTPATH
echo

echo ===== RUN CASE1=====
/hadoop/bin/yarn jar od.jar OutDegree /hw5_data/case1 $OUTPUTPATH"case1"
echo

echo ===== RUN CASE2=====
/hadoop/bin/yarn jar od.jar OutDegree /hw5_data/case2 $OUTPUTPATH"case2"
echo

echo ===== Get Result file =====
/hadoop/bin/hdfs dfs -get $OUTPUTPATH .
echo
```

## 2.1 Explanation of the Map Function

The map function reads each line and uses a `StringTokenizer` to break it into tokens. This is the code in my program:

```
public void map(Object key, Text value, Context context
                ) throws IOException, InterruptedException {
    StringTokenizer itr = new StringTokenizer(value.toString());
    itr.nextToken(); // Skip the first token ("a")
    if (itr.hasMoreTokens()) {
        String sourceNode = itr.nextToken(); // The source node
        node.set(sourceNode);
        // Emit the source node with a count of 1 for each outgoing edge
        context.write(node, one);
    }
}
```

Here's how it processes each line:

- The first token (label) is ignored.
- The function then retrieves the source node and stores it in a `Text` variable.
- For each occurrence of a source node, it emits the source node key with a value of 1, indicating one outgoing edge.

## 2.2 Explanation of the Reduce Function

After the Map phase, we run the reduce function. The `reduce` function combines the out-degree counts for each source node based on the output from the map phase. Below is the code used in the program:

```
public void reduce(Text key, Iterable<IntWritable> values,
                  Context context
                  ) throws IOException, InterruptedException {
    int sum = 0;
    for (IntWritable val : values) {
        sum += val.get(); // Sum up all counts for each node
    }
    result.set(sum);
    context.write(key, result); // Emit the node and its total out-degree
}
```

Here's how the reduce function processes the intermediate data:

- For each unique source node key, the reduce function receives an iterable list of `IntWritable` values, where each value represents one outgoing edge emitted by the map function.
- It initializes a counter, `sum`, to accumulate the out-degree.
- The function iterates over each value in the list and adds it to `sum`.
- After calculating the total count for each source node, the result is set to `sum` and emitted as the final out-degree for that node.

## 3 Results from running the MapReduce

### 3.1 Terminal output

Running the run\_od.sh shell script, we get the following terminal output:

```
===== OutDegree JAVA VERSION =====
===== Compile =====
Note: OutDegree.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.

===== Clear old output files on HDFS =====
Deleted od_res

===== RUN CASE1=====
2024-10-27 11:25:36,620 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager
2024-10-27 11:25:37,062 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/h
2024-10-27 11:25:37,340 INFO input.FileInputFormat: Total input files to process : 1
2024-10-27 11:25:37,398 INFO mapreduce.JobSubmitter: number of splits:1
2024-10-27 11:25:37,521 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1727611706614_071
2024-10-27 11:25:37,521 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-10-27 11:25:37,684 INFO conf.Configuration: resource-types.xml not found
2024-10-27 11:25:37,685 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-10-27 11:25:37,742 INFO impl.YarnClientImpl: Submitted application application_1727611706614_071
2024-10-27 11:25:37,776 INFO mapreduce.Job: The url to track the job: http://intro00:8088/proxy/appli
2024-10-27 11:25:37,776 INFO mapreduce.Job: Running job: job_1727611706614_0710
2024-10-27 11:25:43,890 INFO mapreduce.Job: Job job_1727611706614_0710 running in uber mode : false
2024-10-27 11:25:43,892 INFO mapreduce.Job: map 0% reduce 0%
2024-10-27 11:25:47,976 INFO mapreduce.Job: map 100% reduce 0%
2024-10-27 11:25:53,015 INFO mapreduce.Job: map 100% reduce 100%
2024-10-27 11:25:53,028 INFO mapreduce.Job: Job job_1727611706614_0710 completed successfully
2024-10-27 11:25:53,193 INFO mapreduce.Job: Counters: 54
    File System Counters
        FILE: Number of bytes read=70
        FILE: Number of bytes written=553689
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=986
        HDFS: Number of bytes written=37
        HDFS: Number of read operations=8
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
        HDFS: Number of bytes read erasure-coded=0
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
        Rack-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=2440
        Total time spent by all reduces in occupied slots (ms)=2300
        Total time spent by all map tasks (ms)=2440
        Total time spent by all reduce tasks (ms)=2300
        Total vcore-milliseconds taken by all map tasks=2440
        Total vcore-milliseconds taken by all reduce tasks=2300
        Total megabyte-milliseconds taken by all map tasks=2498560
        Total megabyte-milliseconds taken by all reduce tasks=2355200
    Map-Reduce Framework
```

```

Map input records=100
Map output records=100
Map output bytes=600
Map output materialized bytes=70
Input split bytes=99
Combine input records=100
Combine output records=8
Reduce input groups=8
Reduce shuffle bytes=70
Reduce input records=8
Reduce output records=8
Spilled Records=16
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=151
CPU time spent (ms)=1020
Physical memory (bytes) snapshot=663908352
Virtual memory (bytes) snapshot=5085540352
Total committed heap usage (bytes)=1246232576
Peak Map Physical memory (bytes)=334999552
Peak Map Virtual memory (bytes)=2539573248
Peak Reduce Physical memory (bytes)=328908800
Peak Reduce Virtual memory (bytes)=2545967104
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=887
File Output Format Counters
  Bytes Written=37

```

===== RUN CASE2=====

```

2024-10-27 11:25:55,117 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager
2024-10-27 11:25:55,567 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/h
2024-10-27 11:25:55,846 INFO input.FileInputFormat: Total input files to process : 1
2024-10-27 11:25:55,900 INFO mapreduce.JobSubmitter: number of splits:2
2024-10-27 11:25:56,014 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1727611706614_071
2024-10-27 11:25:56,014 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-10-27 11:25:56,179 INFO conf.Configuration: resource-types.xml not found
2024-10-27 11:25:56,180 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-10-27 11:25:56,243 INFO impl.YarnClientImpl: Submitted application application_1727611706614_071
2024-10-27 11:25:56,277 INFO mapreduce.Job: The url to track the job: http://intro00:8088/proxy/appli
2024-10-27 11:25:56,278 INFO mapreduce.Job: Running job: job_1727611706614_0711
2024-10-27 11:26:02,402 INFO mapreduce.Job: Job job_1727611706614_0711 running in uber mode : false
2024-10-27 11:26:02,404 INFO mapreduce.Job: map 0% reduce 0%
2024-10-27 11:26:12,518 INFO mapreduce.Job: map 50% reduce 0%
2024-10-27 11:26:18,581 INFO mapreduce.Job: map 81% reduce 0%
2024-10-27 11:26:24,618 INFO mapreduce.Job: map 100% reduce 0%
2024-10-27 11:26:25,624 INFO mapreduce.Job: map 100% reduce 100%
2024-10-27 11:26:25,635 INFO mapreduce.Job: Job job_1727611706614_0711 completed successfully
2024-10-27 11:26:25,768 INFO mapreduce.Job: Counters: 56

```

#### File System Counters

FILE: Number of bytes read=19623412  
FILE: Number of bytes written=29200919  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=179691884  
HDFS: Number of bytes written=3802485  
HDFS: Number of read operations=11  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2  
HDFS: Number of bytes read erasure-coded=0

#### Job Counters

Killed map tasks=1  
Launched map tasks=3  
Launched reduce tasks=1  
Data-local map tasks=1  
Rack-local map tasks=2  
Total time spent by all maps in occupied slots (ms)=32400  
Total time spent by all reduces in occupied slots (ms)=10422  
Total time spent by all map tasks (ms)=32400  
Total time spent by all reduce tasks (ms)=10422  
Total vcore-milliseconds taken by all map tasks=32400  
Total vcore-milliseconds taken by all reduce tasks=10422  
Total megabyte-milliseconds taken by all map tasks=33177600  
Total megabyte-milliseconds taken by all reduce tasks=10672128

#### Map-Reduce Framework

Map input records=10000000  
Map output records=10000000  
Map output bytes=105233939  
Map output materialized bytes=8747156  
Input split bytes=198  
Combine input records=10847655  
Combine output records=1528498  
Reduce input groups=409565  
Reduce shuffle bytes=8747156  
Reduce input records=680843  
Reduce output records=409565  
Spilled Records=2209341  
Shuffled Maps =2  
Failed Shuffles=0  
Merged Map outputs=2  
GC time elapsed (ms)=661  
CPU time spent (ms)=32240  
Physical memory (bytes) snapshot=1595547648  
Virtual memory (bytes) snapshot=7648034816  
Total committed heap usage (bytes)=2345664512  
Peak Map Physical memory (bytes)=630136832  
Peak Map Virtual memory (bytes)=2552717312  
Peak Reduce Physical memory (bytes)=365457408  
Peak Reduce Virtual memory (bytes)=2550886400

#### Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0

```
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=179691686
File Output Format Counters
  Bytes Written=3802485

==== Get Result file ====
get: 'od_res/case1/_SUCCESS': File exists
get: 'od_res/case1/part-r-00000': File exists
get: 'od_res/case2/_SUCCESS': File exists
get: 'od_res/case2/part-r-00000': File exists
```

Here we see that the code has not found any errors when running case1 and case2. Still, we will discuss the result of these two cases run by my algorithm



### 3.2 Case 1

When the program is run, we can go to `od_res/case1/part-r-00000`. By opening file we get the following:

```
1 31
2 18
3 13
4 10
6 15
7 3
8 8
9 2
```

This is the correct structure as described by the *ReadMe*, where every number is described once and their instances is recorded as a sum to their right.

### 3.3 Case 2

We can also open `od_res/case2/part-r-00000`. This file has 409565 lines and therefore unsuitable to display in this report. I will, however, show the first and last 5 entries:

```
1      14668
100     372
1000  363
100001 6
100002 1
...
99994 12
999941 1
99995 4
99998 18
99999 5
```

This file can also be found in the code also found in the homework.